

Architecture of LISP Machines

Kshitij Sudan

March 6, 2008

A Short History Lesson ...

Alonzo Church and Stephen Kleene (1930) – λ Calculus

(to cleanly define "computable functions")



John McCarthy (late 60's)

(used λ Calculus to describe the operation of a computing machine to prove theorems about computation)




MIT → *"Knights of the Lambda Calculus"*



MIT AI Lab (~1970's)



Symbolics and LMI



“MacLisp” family Machines

1975	The CONS prototype (MIT)	
1977	The CADR aka MIT Lisp Machine (MIT)	
1980	LM-2 Symbolics Lisp Machine, repackage CADR	LMI Lisp Machine same as CADR
1982	L-Machine - Symbolics 3600, later 3640, 3670	
1983	LMI Lambda	TI Explorer same as LMI Lambda
1984	G-Machine - Symbolics 3650	
1986	LMI K-Machine	
1987	I-Machine , Symbolics XL-400, Macivory I	TI Explorer-II - u-Explorer
1988	Macivory II	
1989	I-Machine , Symbolics XL-1200 , Macivory III	
1990	XL1200, UX-1200	
1991	MacIvory III	
1992	Virtual Lisp Machine (aka Open Genera) I-machine compatible, running on DEC Alpha	

Agenda

- History of LISP machines.
- Semantic Models.
- von Neumann model of computation.
- Programming language to m/c architecture.
- Architectural challenges.
- The SECD abstract machine.
- A brief case study.

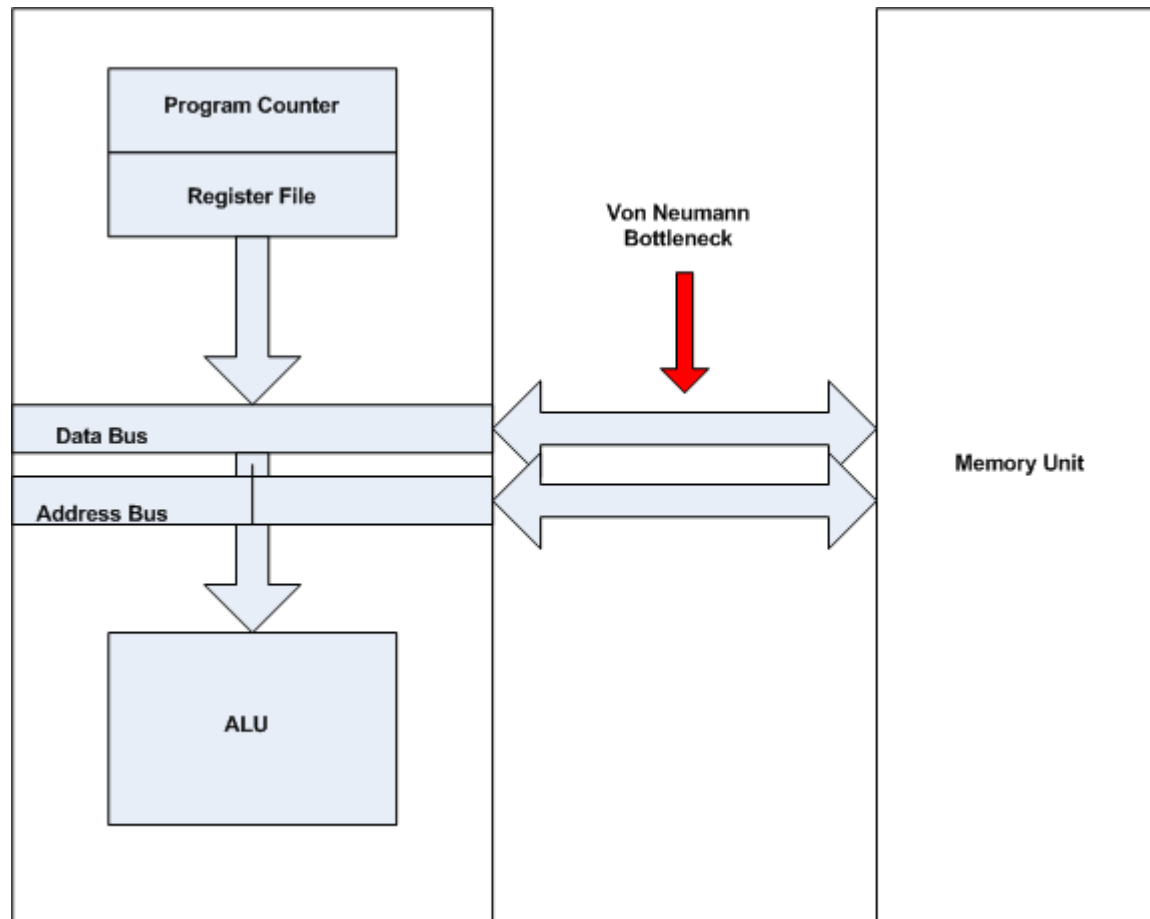
Semantic Models

- The ***semantics*** of a piece of notation is its ultimate meaning.

Imp. for programmer → Imp. for language designer → Imp. for architects

- Three major methods to describe and define semantics of programming languages:
 - ***Interpretive*** : meaning is expressed in terms of some simple abstract m/c.
 - ***Axiomatic*** : where rules describe data values given various objects before and after execution of various language features.
 - ***Denotational*** : syntactic pieces of program are mapped via evaluation functions into the abstract values they denote to humans.

von Neumann Model of Computation



Programming Languages to Machine Architectures

- **Interplay** between h/w (m/c org.) and s/w (compilers, interpreters, and run-time routines) needs to be sustainable for efficient computational structures.
- **Mathematical framework** → Computing models → languages → architecture → real implementations.
- Mathematical framework → ***Abstract m/c*** → **real implementations.**

A short detour ...

- Processing symbols “*was*” touted (circa early 90’s) as future of computations (*obviously hasn’t happened yet!*)
- For processing symbols, declarative languages were put forth as the solution –
 - ***function-based*** and ***logic-based*** languages

So what is the future?

Architectural challenges - I

- Today we talk mostly about LISP machines (functional language m/c's).
- Describe features “*needed*” for *efficient* LISP program execution (RISC can obviously execute LISP).
- Language feature driven architectural hooks – we talk about then briefly.
- Abstract m/c → case studies

Architectural challenges – II

(Architectural support for LISP - I)

- Fast function calls.
 - *call* and *return* instructions with short execution latencies for dynamically bound contexts (latest active value bound to a variable name).
 - *funarg problem*.
- Environment maintenance.
 - *shallow-bound* (linked-list)
 - *deep-bound* (“oblist” == global symbol table)
 - with possible *caching of name-value bindings* (value cache).

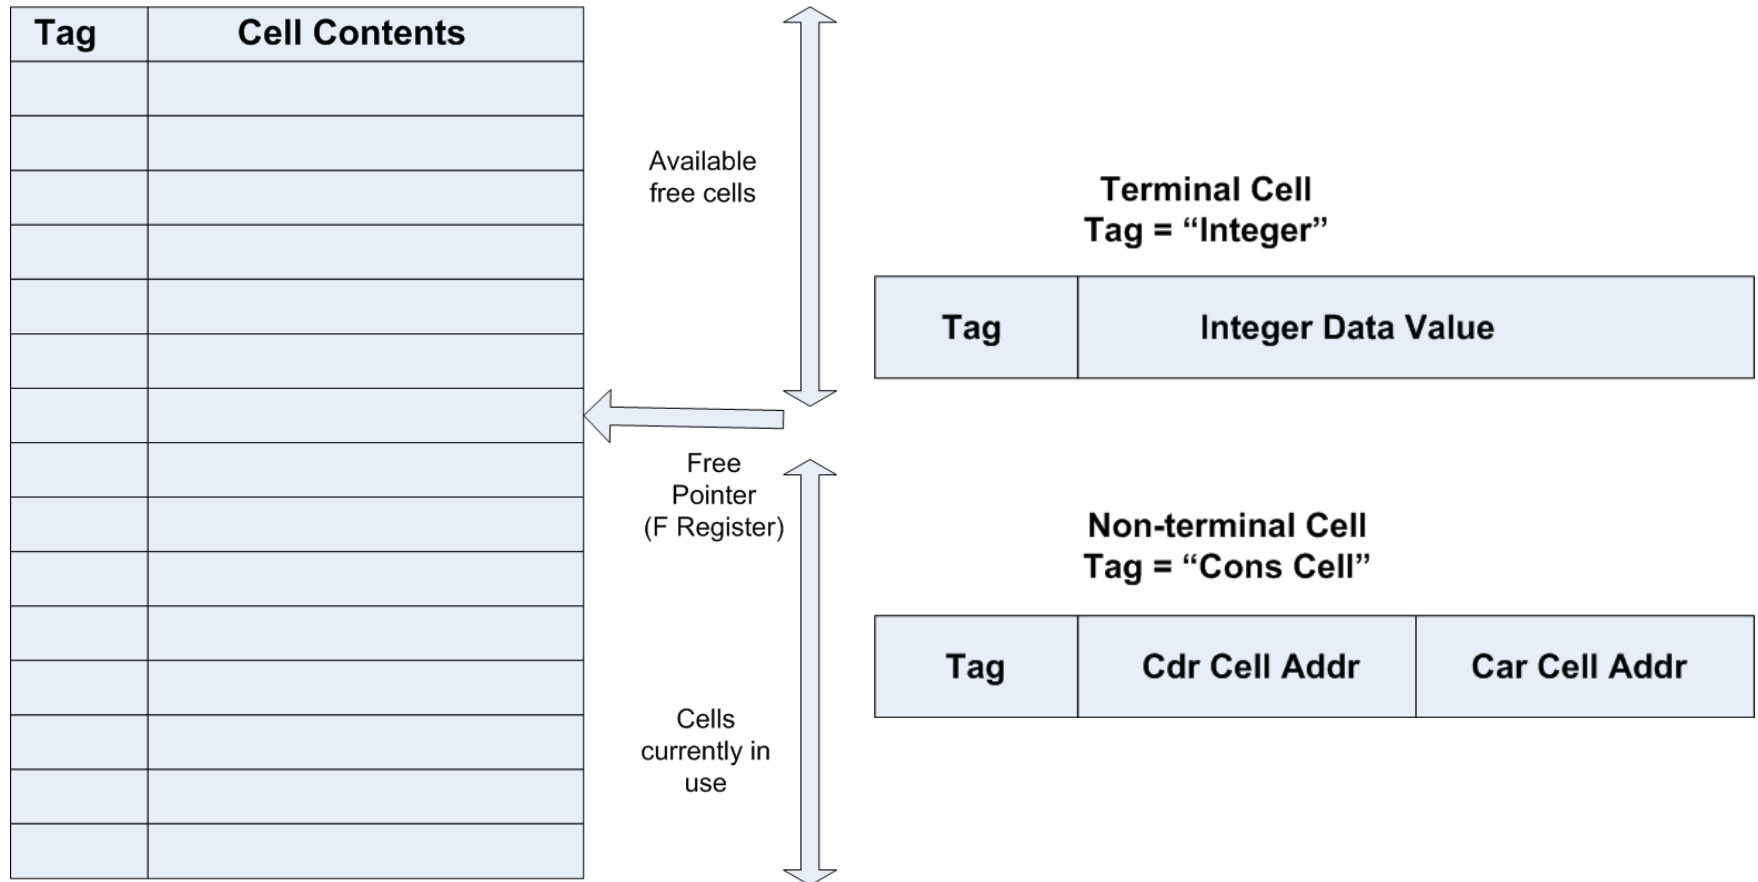
Architectural challenges – III

(Architectural support for LISP - II)

- Efficient list representation.
 - improvements over ***two-pointer list cells***
 - ***Vector-coded*** (represent linear lists as vector of symbols)
 - ***Structure-coded***.
 - each cell has a tag for it's location in the list.
 - associative search leads to fast access.
- Heap maintenance (a.k.a. garbage collection)
 - ***Marking*** (accessible lists “marked”, others reclaimed)
 - ***Reference count*** (count links to the cell, when ==0, reclaim)
 - Generally mix of two schemes used.
- Dynamic type checking.
 - tagged memories and special type-checking h/w

The SECD Abstract Machine

Memory



The SECD Abstract Machine

Basic Data Structures

- Arbitrary s-expressions for computed data.
- List representing programs to be executed.
- Stack's used by programs instructions.
- Value Lists containing arguments for uncompleted function applications.
- Closures to represent unprocessed function applications.

The SECD Abstract Machine

Machine Registers

- **S – Register** (Stack register)
 - Points to a list in memory that's treated as a conventional stack **for built-in functions** (+, -, etc)
 - Objects to be processed are pushed on by **cons'**ing a new cell on top of the current stack and **car** of this points to object's value.
 - S- register after such a push points to the new cell.
 - Unlike conventional stack, this **does not overwrite original inputs.**
 - Cells garbage collected later.

The SECD Abstract Machine

Machine Registers

- **E – Register** (Environment register)
 - Points to current value list of function arguments
 - The list is **referenced by m/c when a value for the argument is needed.**
 - List is **augmented when a new environment for a function is created.**
 - It's **modified when a previously created closure is unpacked** and the pointer from the closure's **cdr** replaces the contents of E-register.
 - Prior value list designated by E is not overwritten.

The SECD Abstract Machine

Machine Registers

- **C – Register** (Control register/pointer)
 - Acts as the program counter and **points to the memory cell that designates through it's car the next instruction to be executed.**
 - The instructions are simple integers specifying desired operation.
 - **Instructions do not have any sub-fields** for registers etc. If additional information is required, it's accessed through from the cells chained through the instruction cell's **cdr**.
 - **“Increment of PC”** takes place by replacement of C registers contents by the contents of the last cell used by the instruction.
 - For **return from completed applications**, new function calls and branches, the C register is replaced by a pointer provided by some other part of the m/c.

The SECD Abstract Machine

Machine Registers

- **D – register** (Dump register)
 - Points to a list in memory called “**dump**”.
 - This data structure **remembers the state of a function** application when a new application in that function body is started.
 - That is done by appending onto dump the 3 new cells which record in their **cars** the value of registers **S**, **E**, and **C**.
 - When the application completes, popping the top of the dump restores those registers. This **is very similar to call-return sequence in conventional m/c for procedure return and activation.**

The SECD Abstract Machine

Basic Instruction Set

- Instruction can be classified into following 6 groups:
 1. Push object values onto the S stack.
 2. Perform built-in function applications on the S stack and return the result to that stack.
 3. Handle the if-then-else special form.
 4. Build, apply and return from closures representing non-recursive function applications.
 5. Extend the above to handle recursive functions.
 6. Handle I/O and machine control.

The CADR machine built at MIT (1984) closely resembles SECD with some non-trivial differences.

Case Study

Concert machine for MultiLISP (1985)

- **MultiLISP**
 - designed as an extension of SCHEME that **permits the programmer to specify parallelism** and then supports the parallelism in h/w “efficiently”.
- **SCHEME + new calls:**
 1. (PCALL F E1 E2 ... En)
 - Permit parallel evaluation of arguments, then evaluate (F E1 E2 ... En)
 2. (DELAY E)
 - Package E in closure.
 3. (TOUCH E)
 - Do not return until E evaluated.
 4. (FUTURE E)
 - Package E in a closure and permit eager evaluation
 5. (REPLACE-xxx E1 E2) [xxx is either CAR or CDR]
 - Replace xxx component of E1 by E2. (**permits controlled modification to storage**)
 6. (REPLACE-xxx-EQ E1 E2 E3)
 - Replace xxx of E1 by E2 iff xxx = E3. (**TEST_AND_SET**)

Case Study

Concert machine for MultiLISP (1985)

- Concert m/c at MIT – 24-way Motorola 68000 based shared memory multiprocessor.
- MultiLISP → **MCODE** (SECD-like ISA) → Interpreted by C interpreter (~ 3000 loc)
- Common gc heap distributed among all processor memories to hold all shared data.
- MCODE programs manage data structure called **tasks** that are accessed by 3 pointers: *program pointer, stack pointer, and environment pointer.*

Case Study

Concert machine for MultiLISP (1985)

- *FUTURE* call creates a new task and leaves it accessible for any free processor. It's environment is that of it's parent expression at it's time creation.
- ***Task queue*** used to maintain schedulable tasks and ***unfair scheduling policy*** used to prevent task explosion.
- GC uses **Banker's algorithm** and spread over all processors with **careful synchronization** to avoid multiple processors trying to evacuate same object at the same time.

Questions?

Thanks for your patience ...